

4-6 Wednesday – 210-GD3

# Special topics in Computer Science

## INT3121 20

Lecturer: Nguyen Thi Ngoc Diep, Ph.D.

Email: [ngocdiep@vnu.edu.vn](mailto:ngocdiep@vnu.edu.vn)

Slide & Code: <https://github.com/chupibk/INT3121-20>

## Image classification with convolutional neural networks

| Week                  | Content   | Class hour | Self-study hour |
|-----------------------|---|------------|-----------------|
| <b>1</b><br>28/8/2019 | Introduction<br>Image classification problem and its applications<br>A toy problem with CIFAR10 | 2          | 1               |
| <b>2</b>              | <b>CNN model architectures and visualization</b>  | <b>2</b>   | <b>1</b>        |
| <b>3</b>              | Training and tuning parameters<br>Automatic parameter learning                                  | 2          | 1               |
| <b>4</b>              | Data augmentation<br>Data generator   | 2          | 2-6             |
| <b>5</b>              | Transfer learning   | 2          | 2-6             |
| <b>6</b>              | Multi-output image classification   | 2          | 2-6             |
| <b>7</b>              | Building a training dataset<br>How to write a report  | 1          | 2-6             |
| <b>8, 9, 10, 11</b>   | Seminar: Bag of tricks with CNN<br>(as mid-term tests)  | 1          | 2-6             |
| <b>12, 13, 14</b>     | Final project presentations   | 1-3        | 2-6             |
| <b>15</b>             | Class summarization   | 1          | open            |


## Week 1: Homework checklist

- Install all necessary environments?
  - Keras, tensorflow, python 3, jupyter notebook, numpy, sklearn
- Re-run all the scripts?
- Obtain the label names of CIFAR10 classes?
- Given an image, predict the label of the image?

→ Check lecture code for answers 😊

## Layers in Neural Networks

## A sample net with fully connected layers



Visualization tool: <http://alexlenail.me/NN-SVG/index.html>

Image credit: <https://www.learnopencv.com/understanding-activation-functions-in-deep-learning/>

## Code it with Keras

```
from keras.models import Sequential
from keras.layers import *

model = Sequential()
model.add(Dense(4, input_shape=(3,)))
model.add(Dense(5))
model.add(Dense(2))
model.summary()
```

| Layer (type)            | Output Shape | Param # |
|-------------------------|--------------|---------|
| dense_4 (Dense)         | (None, 4)    | 16      |
| dense_5 (Dense)         | (None, 5)    | 25      |
| dense_6 (Dense)         | (None, 2)    | 12      |
| Total params: 53        |              |         |
| Trainable params: 53    |              |         |
| Non-trainable params: 0 |              |         |

Where does this come from?

# Convolution Layers

## LeNet example





Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Reference: LeCun, Yann, et al. "Gradient-based learning applied to document recognition." *Proceedings of the IEEE* 86.11 (1998): 2278-2324.  
<http://yann.lecun.com/exdb/publis/pdf/lecun-98.pdf>

## Spatial filtering



Gonzalez et al.

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image



|   |  |  |
|---|--|--|
| 4 |  |  |
|   |  |  |
|   |  |  |

Convolved  
Feature

## Spatial filtering as feature extraction

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * I$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * I$$




Horizontal changes



Vertical changes


## Sobel in both directions



[https://en.wikipedia.org/wiki/Sobel\\_operator](https://en.wikipedia.org/wiki/Sobel_operator)


Kernel = function  
Convolution = operation to combine 2 functions

Input



```
arr = [
  [0, 0, 1, 0, 0, 0],
  [0, 1, 1, 0, 0, 0],
  [1, 0, 0, 0, 0, 1],
  [0, 0, 0, 0, 1, 0],
  [0, 0, 0, 1, 0, 1]
]
```

Template



```
template = [
  [0, 0, 1],
  [0, 1, 1],
  [1, 0, 0]
]
```

Correlation operation




Convolution operation



## Convolution layers are about kernels and their behaviors


- Having “gap” in kernels?
  - Traditional vs Dilated (a.k.a Atrous)
- Dealing with pixels at border?
  - Padding
- Moving step across images?
  - Stride
- Accelerating computation?
  - Separable
- Upsampling data?
  - Deconvolution (aka, Transposed convolution)

## Stride & Padding




Credit: [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

# Convolution for multichannel input



Credit: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>


# Dilated convolution



Credit: [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)



## Upsampling Deconvolution (Transposed convolution)



Credit: [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)


## Important keywords to remember

- Filter
- Kernel
- Stride
- Pad
- Receptive field
- Downsampling
- Upsampling
- Feature map

## Types of filters

- (Traditional) Convolution
  - Too many parameters
- Dilated or Atrous convolution
  - Large receptive field
  - Less parameters
- Separable convolution
  - Even lesser parameters
- Deconvolution or transposed convolution
  - Low resolution -> high resolution
- Pooling -> not a convolution layer, but...

## Pooling layers



## Parameters of a Convolution layer

**Summary.** To summarize, the Conv Layer:

- Accepts a volume of size  $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
  - Number of filters  $K$ ,
  - their spatial extent  $F$ ,
  - the stride  $S$ ,
  - the amount of zero padding  $P$ .
- Produces a volume of size  $W_2 \times H_2 \times D_2$  where:
  - $W_2 = (W_1 - F + 2P)/S + 1$
  - $H_2 = (H_1 - F + 2P)/S + 1$  (i.e. width and height are computed equally by symmetry)
  - $D_2 = K$
- With parameter sharing, it introduces  $F \cdot F \cdot D_1$  weights per filter, for a total of  $(F \cdot F \cdot D_1) \cdot K$  weights and  $K$  biases.
- In the output volume, the  $d$ -th depth slice (of size  $W_2 \times H_2$ ) is the result of performing a valid convolution of the  $d$ -th filter over the input volume with a stride of  $S$ , and then offset by  $d$ -th bias.

A common setting of the hyperparameters is  $F = 3, S = 1, P = 1$ . However, there are common conventions and rules of thumb that motivate these hyperparameters. See the [ConvNet architectures](#) section below.

Reference: <http://cs231n.github.io/convolutional-networks/>

## CNN architectures

## Some CNN architectures




Figure 2: **Top1 vs. operations, size  $\propto$  parameters.**

(Operations required for a single forward pass)

Reference: Canziani, Alfredo, Adam Paszke, and Eugenio Culurciello. "An analysis of deep neural network models for practical applications." *arXiv preprint arXiv:1605.07678* (2016).

## Architectures

- LeNet-5 – LeCun et al., 1998
- AlexNet – Krizhevsky et al., 2012
- VGGNet – Simonyan et al., 2014
- GoogleNet/Inception – Szegedy et al., 2014
- Resnet – Kaiming et al., 2015

## LeNet-5




Image credit: LeCun

## AlexNet




Image credit:...

## VGG16




Image credit: ...

## GoogLeNet/Inception




Image credit: ...

## ResNet with “skip connections”




Image credit: ...