

Special topics in Computer Science

INT3121 20

Lecturer: Nguyen Thi Ngoc Diep, Ph.D.

Email: ngocdiep@vnu.edu.vn

Slide & Code: <https://github.com/chupibk/INT3121-20>

Week 4 recall

- Generator: generate object data one at a time and only when asked
- Training with generator:
 - `model.fit_generator()`
 - Make sure the generator input return values whenever called
 - Data are flowed non-stop
- Augmentation:
 - Operators: blur, sharpen, edge detection, adding noise, invert, pad, crop, flip, resize...
- Generator + augmentation together:
 - `keras.preprocessing.image.ImageDataGenerator()`
 - `.flow(x,y)`
 - `.flow_from_directory(directory)`

Registration (deadline: Oct 2, 2019 23:59)

- Link: <https://forms.gle/BgAWdsCjHgD1nwoA6>
- Group name
- Leader email
- Member IDs
- Chosen dataset
- Desired presentation date
- References: → submit before 1 week
 - Research papers (pdfs)
 - Links
 - Books (pdfs if possible)
- Presentation materials & code → submit before 1 day

Mid-term != final projects

Objectives

- | | |
|---|---|
| <ul style="list-style-type: none"> • Mid-term projects <ul style="list-style-type: none"> • Skills to do survey and comparative studies • Compare peer-reviewed methods on benchmarked datasets <ul style="list-style-type: none"> • <u>Not</u>: I make this & that → how cool I am | <ul style="list-style-type: none"> • Final projects <ul style="list-style-type: none"> • Skills to implement and apply to new problems • Use existing techniques to self-created datasets |
|---|---|

Don't try too hard to be cool ☺
(Let's learn how cool others are)

Now you can ignore others :D

Marking policy in mid-term projects

- | | | |
|---|-------|--|
| • Comparisons of methods/techniques → 30% | → | If failed to do a thorough comparative study, the mark can't exceed 6.0! |
| • More is better | | |
| • Implementation | → 20% | |
| • Analysis of the results | → 20% | |
| • Presentation | → 10% | |
| • Audiences' comments | → 20% | |


Week 5: Transfer learning

Image classification with convolutional neural networks

Week	Content	Class hour	Self-study hour
1 28/8/2019	Introduction Image classification problem and its applications A toy problem with CIFAR10	2	1
2 (4/9/2019)	CNN model architectures and visualization	2	1
3 (11/9/2019)	Training and tuning parameters Automatic parameter learning	2	1
4 (18/9/2019)	Data augmentation Data generator	2	2-6
5 (25/9/2019)	Transfer learning	2	2-6
6 (2/10/2019)	Multi-output image classification	2	2-6
7 (9/10/2019)	Building a training dataset How to write a report	1	2-6
8, 9, 10, 11	Seminar: Bag of tricks with CNN (as mid-term tests)	1	2-6
12, 13, 14	Final project presentations	1-3	2-6
15	Class summarization	1	open

Transfer learning

- Transfer learning = a technique where a model trained on one task is exploited in another task
- Weights of a Convolutional Network are often random initialized
 - Starting of optimizer to find global minimum for loss function
- Transfer learning -> a way to reuse weights




Visualization credit: <https://github.com/Jaewan-Yun/optimizer-visualization>

Three major Transfer learning scenarios

1. ConvNet as fixed feature extractor
 - Remove (some) last fully-connected layer then run forward pass to extract features
2. Fine-tuning the ConvNet
 - Do backpropagation on some last layers only
3. Pretrained models
 - Do backpropagation on the whole model but with pretrained weights

Reference: <http://cs231n.github.io/transfer-learning/>

ConvNet as fixed feature extractor



What to do with features?

- Image classification using other machine learning techniques
 - E.g., SVM
- Image retrieval
 - Similarity-based
- Image clustering
 - Group images into clusters

Assumption of Fine-tuning

- Earlier features of a ConvNet contain more generic features (e.g., edge, texture, color blob)
 - → should be useful to many task
- Later layers of the ConvNet are more specific to the task

When & how to use which transfer learning?

		Size of new dataset	
		Small	Large
Type of new dataset	Similar	<ul style="list-style-type: none"> - Best: ConvNet as fixed feature extractor, then train with a linear classifier - Fine-tune may cause overfitting 	<ul style="list-style-type: none"> - fine-tune is okay
	Very different	<ul style="list-style-type: none"> - Best: Extract earlier features in a ConvNet, then train with a linear classifier 	<ul style="list-style-type: none"> - Initialize the network with weights from a pretrained model

Size of input image?

VGG16 with original
224x224 input

Layer (type)	Output Shape	Param #
input_26 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

100x200 input

Layer (type)	Output Shape	Param #
input_25 (InputLayer)	(None, 100, 200, 3)	0
block1_conv1 (Conv2D)	(None, 100, 200, 64)	1792
block1_conv2 (Conv2D)	(None, 100, 200, 64)	36928
block1_pool (MaxPooling2D)	(None, 50, 100, 64)	0
block2_conv1 (Conv2D)	(None, 50, 100, 128)	73856
block2_conv2 (Conv2D)	(None, 50, 100, 128)	147584
block2_pool (MaxPooling2D)	(None, 25, 50, 128)	0
block3_conv1 (Conv2D)	(None, 25, 50, 256)	295168
block3_conv2 (Conv2D)	(None, 25, 50, 256)	590080
block3_conv3 (Conv2D)	(None, 25, 50, 256)	590080
block3_pool (MaxPooling2D)	(None, 12, 25, 256)	0
block4_conv1 (Conv2D)	(None, 12, 25, 512)	1180160
block4_conv2 (Conv2D)	(None, 12, 25, 512)	2359808
block4_conv3 (Conv2D)	(None, 12, 25, 512)	2359808
block4_pool (MaxPooling2D)	(None, 6, 12, 512)	0
block5_conv1 (Conv2D)	(None, 6, 12, 512)	2359808
block5_conv2 (Conv2D)	(None, 6, 12, 512)	2359808
block5_conv3 (Conv2D)	(None, 6, 12, 512)	2359808
block5_pool (MaxPooling2D)	(None, 3, 6, 512)	0

32x32 input

Layer (type)	Output Shape	Param #
input_23 (InputLayer)	(None, 32, 32, 3)	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block4_pool (MaxPooling2D)	(None, 2, 2, 512)	0
block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0

If continue, run out of data